

Class 8 Mini Project

Emma Bell (A19247017)

Table of contents

Background	1
Principle Component Analysis (PCA)	4
Variance explained	8
Communicating PCA results	11
Hierarchical Clustering	12
Combining methods	13
Prediction	16

Background

In today's class we will be employing all the R techniques for dada analysis that we have learned thus far - including the machine learning methods of clustering and PCA - to analyse real breast cancer biopsy data.

Values in this data set describe characteristics of the cell nuclei present in digitized images of a fine needle aspiration (FNA) of a breast mass.

FNA is a type of biopsy procedure where a very thin needle is inserted into an area of abnormal tissue or cells with a guide of CT scan or ultrasound monitors. The collected sample is then transferred to a pathologist to study it under a microscope and examine whether cells in the biopsy are normal or not.

Features measured from the digitized images include:

radius: mean of distances from nucleus center to points on the perimeter; texture: a measure of nucleus roughness taken from the standard deviation of gray-scale values; perimeter: total boundary length of the nucleus, area: total area of the nucleus; smoothness: local variation in radius lengths, i.e. how "bumpy" the edge is; compactness: measures how circular vs. irregular the shape is; concavity: how deeply indented, and symmetry: how symmetric the nucleus is.

```
wisc.df <- read.csv("WisconsinCancer.csv", row.names=1)
```

wee peak at the data

```
head(wisc.df, 3)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.8	1001
842517	M	20.57	17.77	132.9	1326
84300903	M	19.69	21.25	130.0	1203
	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean	
842302	0.11840	0.27760	0.3001	0.14710	
842517	0.08474	0.07864	0.0869	0.07017	
84300903	0.10960	0.15990	0.1974	0.12790	
	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se
842302	0.2419		0.07871	1.0950	0.9053
842517	0.1812		0.05667	0.5435	0.7339
84300903	0.2069		0.05999	0.7456	0.7869
	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se
842302	153.40	0.006399	0.04904	0.05373	0.01587
842517	74.08	0.005225	0.01308	0.01860	0.01340
84300903	94.03	0.006150	0.04006	0.03832	0.02058
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst	
842302	0.03003	0.006193	25.38	17.33	
842517	0.01389	0.003532	24.99	23.41	
84300903	0.02250	0.004571	23.57	25.53	
	perimeter_worst	area_worst	smoothness_worst	compactness_worst	
842302	184.6	2019	0.1622	0.6656	
842517	158.8	1956	0.1238	0.1866	
84300903	152.5	1709	0.1444	0.4245	
	concavity_worst	concave.points_worst	symmetry_worst		
842302	0.7119	0.2654	0.4601		
842517	0.2416	0.1860	0.2750		
84300903	0.4504	0.2430	0.3613		
	fractal_dimension_worst				
842302	0.11890				
842517	0.08902				
84300903	0.08758				

Q. How many observations are in this dataset

```
nrow(wisc.df)
```

```
[1] 569
```

Q. How many observations have a malignant diagnosis?

```
sum(wisc.df$diagnosis == "M")
```

```
[1] 212
```

or

```
table(wisc.df$diagnosis)
```

	B	M
357	212	

Q. How many variables/features in the data are suffixed with _mean?

```
colnames(wisc.df)
```

```
[1] "diagnosis"                  "radius_mean"  
[3] "texture_mean"              "perimeter_mean"  
[5] "area_mean"                 "smoothness_mean"  
[7] "compactness_mean"           "concavity_mean"  
[9] "concave.points_mean"       "symmetry_mean"  
[11] "fractal_dimension_mean"    "radius_se"  
[13] "texture_se"                "perimeter_se"  
[15] "area_se"                   "smoothness_se"  
[17] "compactness_se"             "concavity_se"  
[19] "concave.points_se"         "symmetry_se"  
[21] "fractal_dimension_se"      "radius_worst"  
[23] "texture_worst"              "perimeter_worst"  
[25] "area_worst"                 "smoothness_worst"  
[27] "compactness_worst"          "concavity_worst"  
[29] "concave.points_worst"       "symmetry_worst"  
[31] "fractal_dimension_worst"
```

```
length(grep("_mean", colnames(wisc.df)))
```

```
[1] 10
```

We need to remove the `diagnosis` column before we do any further analysis of this dataset - we don't want to pass this to PCA etc. We will save it as a separate wee vector that we can use later to compare out findings tho those of experts.

```
diagnosis <- wisc.df$diagnosis
wisc.data <- wisc.df[,-1]
```

Principle Component Analysis (PCA)

The main function in base R is called `prcomp()` we will use the optional `scale=TRUE` here as the data columns/features/dimenstions are on very different scales in the orginial data set.


```
wisc.pr <- prcomp(wisc.data, scale=T)
```

```
attributes(wisc.pr)
```

```
$names
[1] "sdev"      "rotation"   "center"    "scale"     "x"
```

```
$class
[1] "prcomp"
```

```
library(ggplot2)
ggplot(wisc.pr$x) +
  aes(PC1, PC2, col=diagnosis) +
  geom_point()
```



```
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Q4. From your results, what proportion of the original variance is captured by the first principal component (PC1)?

0.4427


Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

PC3 onwards

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

PC7 onwards

```
biplot(wisc.pr)
```




Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

Very chaotic! Very difficult to understand. So lets generate a more standard scatter plot of each observation along principal components 1 and 2.

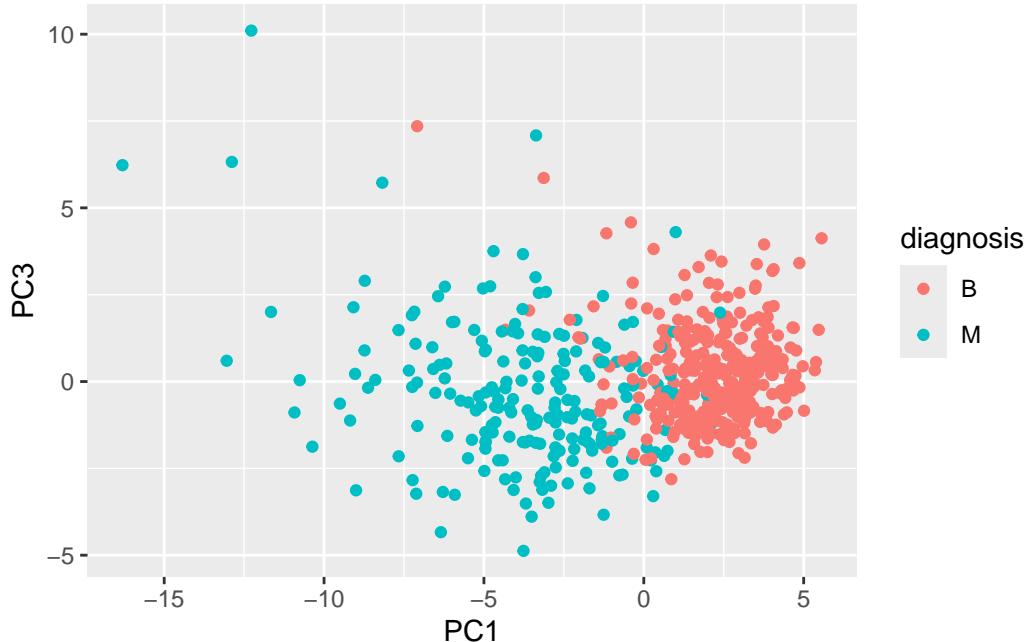
this is the graph from above:

```
library(ggplot2)
ggplot(wisc.pr$x) +
  aes(PC1, PC2, col=diagnosis) +
  geom_point()
```



Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

```
library(ggplot2)
ggplot(wisc.pr$x) +
  aes(PC1, PC3, col=diagnosis) +
  geom_point()
```



This plot is relatively different to PC1 vs PC2, but essentially shows the same distribution of B vs M.

Variance explained

A scree plot shows how much variance each PC captures. We typically look for an “elbow” — a point where adding more PCs gives diminishing returns. This can help us decide how many PCs to consider for further analysis. (Spoiler: some real data sets don’t have a perfect elbow, so folks will often use a threshold like 70% or 90% cumulative variance instead.)

```
pr.var <- wisc.pr$sdev^2
head(pr.var)
```


```
[1] 13.281608 5.691355 2.817949 1.980640 1.648731 1.207357
```

Calculate the variance explained by each principal component by dividing by the total variance explained of all principal components. Assign this to a variable called pve and create a plot of variance explained for each principal component.

```
pve <- (c(pr.var))/100  
head(pve)
```


```
[1] 0.13281608 0.05691355 0.02817949 0.01980640 0.01648731 0.01207357
```

```
plot(c(1,pve), xlab="Principal Component", ylab = "Proportion of Variance Explained",  
      ylim = c(0, 1), type = "o")
```



Or an alternative barplot:

```
barplot(pve, ylab = "Percent of Variance Explained",  
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)  
axis(2, at=pve, labels=round(pve,2)*100 )
```



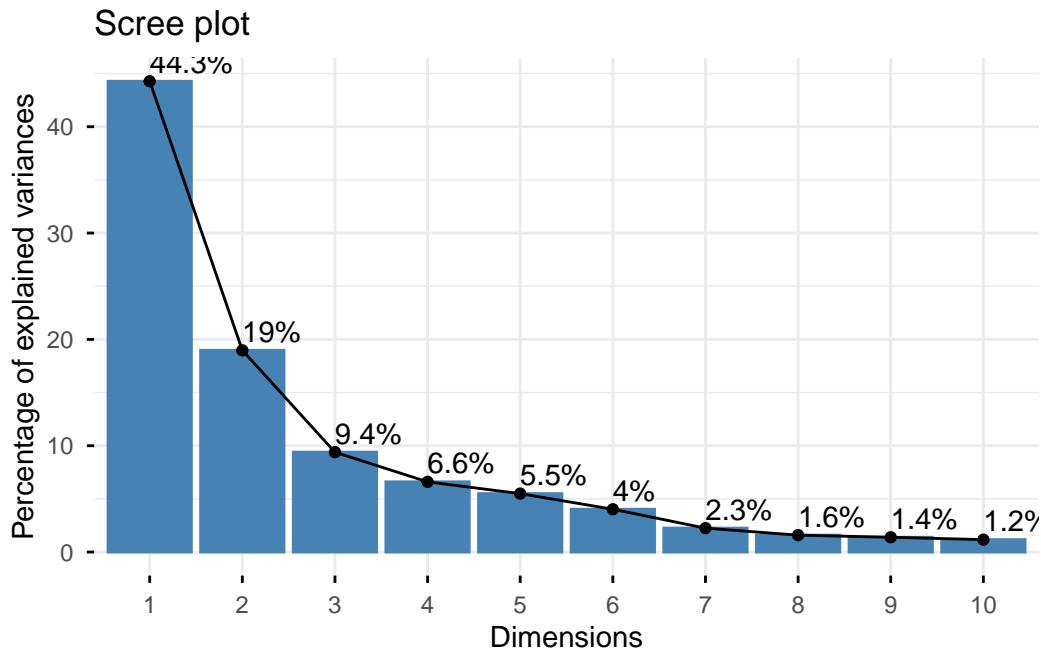
OPTIONAL: There are quite a few CRAN packages that are helpful for PCA. This includes the factoextra package. Feel free to explore this package. For example:

```
library(factoextra)
```

Welcome! Want to learn more? See two factoextra-related books at <https://goo.gl/ve3WBa>

```
fviz_eig(wisc.pr, addlabels = TRUE)
```

Warning in geom_bar(stat = "identity", fill = barfill, color = barcolor, :
Ignoring empty aesthetic: `width`.



Communicating PCA results

In this section we will check your understanding of the PCA results, in particular the “loadings” and “variance explained”.

The loading vector (`wisc.pr$rotation`) tells us which original measurements contribute most to each PC.

A large PC1 loading value (positive or negative) for one of the 30 original measurements (i.e. features/columns we started with), for example, would suggest that this feature is an important driver of the variation we see in the score plot and thus helpful for distinguishing “M” from “B” samples. Let’s check which features matter most to PC1.

Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[1]`) for the feature `concave.points_mean`? This tells us how much this original feature contributes to the first PC. Are there any features with larger contributions than this one?

```
wisc.pr$rotation[8,1]
```

```
[1] -0.2608538
```

```
wisc.pr$rotation[,1]
```

radius_mean	texture_mean	perimeter_mean
-0.21890244	-0.10372458	-0.22753729
area_mean	smoothness_mean	compactness_mean
-0.22099499	-0.14258969	-0.23928535
concavity_mean	concave.points_mean	symmetry_mean
-0.25840048	-0.26085376	-0.13816696
fractal_dimension_mean	radius_se	texture_se
-0.06436335	-0.20597878	-0.01742803
perimeter_se	area_se	smoothness_se
-0.21132592	-0.20286964	-0.01453145
compactness_se	concavity_se	concave.points_se
-0.17039345	-0.15358979	-0.18341740
symmetry_se	fractal_dimension_se	radius_worst
-0.04249842	-0.10256832	-0.22799663
texture_worst	perimeter_worst	area_worst
-0.10446933	-0.23663968	-0.22487053
smoothness_worst	compactness_worst	concavity_worst
-0.12795256	-0.21009588	-0.22876753
concave.points_worst	symmetry_worst	fractal_dimension_worst
-0.25088597	-0.12290456	-0.13178394

so, no, there are no features with larger contributions.

Hierarchical Clustering

The goal of this section is to do hierarchical clustering of the original data to see if there is any obvious grouping into malignant and benign clusters.

Recall from class that hierarchical clustering does not assume in advance the number of natural groups that exist in the data (unlike K-means clustering).

As part of the preparation for hierarchical clustering, the distance between all pairs of observations needs to be calculated. This “distance matrix” will be the input for the `hclust()` function.

One of the optional arguments to `hclust()` allows you to pick different ways (a.k.a. “methods”) to link clusters together, with single, complete, and average being the most common “linkage methods”. You can explore the effects of these different methods in this section.

```


data.scaled <- scale(wisc.data)

data.dist <- dist(data.scaled)

wisc.hclust <- hclust(data.dist, method= "ward.D2")

plot(wisc.hclust)
abline(wisc.hclust, col="red", lty=2)

```




Q12. Which method gives your favorite results for the same data.dist dataset?
Explain your reasoning.

I like `method=ward.D2` as it seems to give a clearer, less chaotic dendrogram.

Combining methods

The idea here is that I can take my new variables (the PCs `wisc.pr$x`) that are better descriptors of the data-set than the original features (ie the 30 columns in `wisc.data`) and use these as a basis for clustering.

```
pc.dist <- dist(wisc.pr$x[, 1:3])
wisc.pr.hclust <- hclust(pc.dist, method= "ward.D2")
plot(wisc.pr.hclust)
```



```
pc.dist
hclust (*, "ward.D2")
```

The next section covers Q13-14.

```
grps <- cutree(wisc.pr.hclust, k=2)
table (grps)
```

```
grps
 1   2
203 366
```

I can now run `table()` with both my clustering `grps` and the expert `diagnosis`

```
table(grps, diagnosis)
```

		diagnosis	
grps	B	M	
1	24	179	
2	333	33	

```
wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k=2)
```

Our cluster “1” has 179 “M” diagnosis Our cluster “2” has 333 “B” diagnosis So cluster “1” is malignant and “2” is benign.

179 TP (true positive) 24 FP 333 TN 33 FN

Sensitivity refers to a test’s ability to correctly detect ill patients who do have the condition. In our example here the sensitivity is the total number of samples in the cluster identified as predominantly malignant (cancerous) divided by the total number of known malignant samples. In other words: $TP/(TP+FN)$.

Q14. How well do the hierarchical clustering models you created in the previous sections (i.e. without first doing PCA) do in terms of separating the diagnoses? Again, use the table() function to compare the output of each model (wisc.hclust.clusters and wisc.pr.hclust.clusters) with the vector containing the actual diagnoses.

```
table(wisc.pr.hclust.clusters, diagnosis)
```

		diagnosis	
		B	M
wisc.pr.hclust.clusters	1	24	179
	2	333	33

It does better because it shows the false positives/false negatives that we can now use.

Sensitivity: $TP/(TP+FN)$

```
179/(179+33)
```

```
[1] 0.8443396
```

sensitivity of 1 would be perfect.

Specificity relates to a test’s ability to correctly reject healthy patients without a condition. In our example specificity is the proportion of benign (not cancerous) samples in the cluster identified as predominantly benign that are known to be benign. In other words: $TN/(TN+FP)$.

Specificity: $TN/(TN+FP)$.

```
333/(333+24)
```

```
[1] 0.9327731
```


Prediction

We can see our PCA model prediction

```
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc
```

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
[1,]	2.576616	-3.135913	1.3990492	-0.7631950	2.781648	-0.8150185	-0.3959098
[2,]	-4.754928	-3.009033	-0.1660946	-0.6052952	-1.140698	-1.2189945	0.8193031
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
[1,]	-0.2307350	0.1029569	-0.9272861	0.3411457	0.375921	0.1610764	1.187882
[2,]	-0.3307423	0.5281896	-0.4855301	0.7173233	-1.185917	0.5893856	0.303029
	PC15	PC16	PC17	PC18	PC19	PC20	
[1,]	0.3216974	-0.1743616	-0.07875393	-0.11207028	-0.08802955	-0.2495216	
[2,]	0.1299153	0.1448061	-0.40509706	0.06565549	0.25591230	-0.4289500	
	PC21	PC22	PC23	PC24	PC25	PC26	
[1,]	0.1228233	0.09358453	0.08347651	0.1223396	0.02124121	0.078884581	
[2,]	-0.1224776	0.01732146	0.06316631	-0.2338618	-0.20755948	-0.009833238	
	PC27	PC28	PC29	PC30			
[1,]	0.220199544	-0.02946023	-0.015620933	0.005269029			
[2,]	-0.001134152	0.09638361	0.002795349	-0.019015820			

```
plot(wisc.pr$x[,1:2], col=grps)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```



Q16. Which of these new patients should we prioritize for follow up based on your results?

We need to prioritise group 2.